

Discussion 11

Graph Algorithms

Objectives

- Shortest path algorithms
 - Dijkstra's Algorithm
 - Bellman-Ford Algorithm
- Minimum spanning tree
 - Prim's Algorithm
 - Kruskal's Algorithm

Shortest Path Algorithms

Overview of Shortest Path Algorithms

1. Single Source Shortest Path Problem

- a. Minimum weight path from a single node to any other node

2. All Pairs Shortest Path

- a. Minimum weight path from **any** node to any other node
- b. Floyd-Warshall algorithm: $O(|V|^3)$
 - i. can have negative weight edges
 - ii. can detect negative weight cycles

Single Source Shortest Path Algorithms

Shortest path from a single source to all vertices

1. **BFS:** $O(|V| + |E|)$
a. unweighted graph <- shortest path in steps
2. **Dijkstra's algorithm:** $O(|V|^2)$, or if implemented by heap $O((|V| + |E|) * \log(|V|))$
a. weighted graph
b. no negative weight edge <- shortest path in weight
3. **Bellman-Ford algorithm:** $O(|V| * |E|)$
a. can have negative weight edges
b. no negative cycles (detects them) <- shortest path in weight

Dijkstra's algorithm

- Determines shortest path to each vertex from a **single source**
- Works on:
 - Weighted graphs
 - **That DON'T have negative weights**
 - Directed graphs OR undirected graphs
- A **greedy** algorithm
 - For each iteration, chooses the **next** vertex to visit by:
 - which **unvisited** vertex has the **minimum** distance to the **source** vertex.
- For each iteration:
 - **Relaxation** principle applied to all the edges that originate from the **unvisited** vertex that has the **minimum** distance to the **source** vertex.
- Time complexity
 - $O(|V|^2)$, or if implemented by heap and adjacency list $O((|V| + |E|) * \log(|V|))$


Edge Relaxation Principle

- Considers an edge and whether the path through vertex u to vertex v is better than any previously found path to vertex v .
- If so, it updates a distance map for key v to be the distance from the source to u plus the distance from u to v .
- Boils down to:

For the edge from the vertex u to the vertex v ,
if $d[u] + w(u, v) < d[v]$ is satisfied, update $d[v]$ to
 $d[u] + w(u, v)$


- The algorithms for the shortest paths problem solve the problem by repeatedly using the edge relaxation in a certain order of relaxation.

Dijkstra's algorithm




- We will cover the Dijkstra's algorithm variant that maintains three maps. For all of these maps, each vertex is a key.
 - **V**: A map that tracks **whether a vertex has been visited**
 - A boolean value is used to indicate whether a vertex has been visited
 - **dist**: A map that **stores all the distances from the source node**
 - For all but source, distance = infinity to start
 - Source vertex distance = 0 (selected first)
 - **P**: A **predecessor map**
 - The value is the last vertex passed through to arrive at the key vertex
 - All values = null to start
 - Used to determine the shortest paths
 - *Ignore if you only want the shortest distance*

Dijkstra's algorithm: Initialization




Key (k)	V[k]	dist[k]	p[k]
s	F	0	null
t	F	∞	null
y	F	∞	null
x	F	∞	null
z	F	∞	null

1st Iteration: s is selected




Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	F	∞	null
y	F	∞	null
x	F	∞	null
z	F	∞	null

1st Iteration: the distances to t and y are relaxed



Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	F	10	s
y	F	5	s
x	F	∞	null
z	F	∞	null

2nd iteration: y is visited




Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	F	10	s
y	T	5	s
x	F	∞	null
z	F	∞	null

Note: Whenever we select a node N to visit, the current distance to N will already be minimized (this is why the algorithm can go node by node and why it doesn't work for negative weights).


Notice that any path to y has to either go through y or t, so 5 is the shortest possible path to y.

2nd iteration: the distances to t, x, and z are relaxed




Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	F	8	y
y	T	5	s
x	F	14	y
z	F	7	y

3rd iteration: z is visited




Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	F	8	y
y	T	5	s
x	F	14	y
z	T	7	y

3rd iteration: the distance to x is relaxed




Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	F	8	y
y	T	5	s
x	F	13	z
z	T	7	y

4th iteration: t is visited




Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	T	8	y
y	T	5	s
x	F	13	z
z	T	7	y

4th iteration: the distance to x is relaxed




Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	T	8	y
y	T	5	s
x	F	9	t
z	T	7	y

5th and final iteration: x is visited



Key (k)	V[k]	dist[k]	p[k]
s	T	0	null
t	T	8	y
y	T	5	s
x	T	9	t
z	T	7	y

Dijkstra's algorithm



Initialize V with all vertices marked as unvisited,
 $dist$ with the source vertex distance value set to 0 and all others set to infinity, and
 P with all values set to null.

While V has an unvisited vertex:

- Find an unvisited vertex u that has the minimum distance value in $dist$.
 - Note: This step is where an optimization can be done using a heap
- Mark that u has been visited in V .
- For each of u 's edges to a vertex y :
 - If y is not visited in V ,
 - **Relax** the distance in the table for y if necessary.
 - i.e. Select the $\min(\text{current value for } y \text{ in } dist, \text{new distance for } y \text{ through } u)$
 - If the distance for y was relaxed, set y 's predecessor value to u in P .

Poll Question #1:

Which of the following statements is not correct about Dijkstra's algorithm?

- A. Dijkstra's algorithm uses a greedy strategy.
- B. The input graph can have negative weight edges.
- C. Dijkstra's algorithm is a single source shortest path algorithm.
- D. The input graph can be undirected or directed.

Poll Question #1:

Which of the following statements is not correct about Dijkstra's algorithm?

- A. Dijkstra's algorithm uses a greedy strategy.
- B. The input graph can have negative weight edges.
- C. Dijkstra's algorithm is a single source shortest path algorithm.
- D. The input graph can be undirected or directed.

Poll Question #2:

Given the following table of vertices and their predecessors from Dijkstra's algorithm, what is the shortest path between s and x?

- A. s, y, t, x
- B. s, y, z, x
- C. s, t, x
- D. s, y, x

Key (k)	d[k]	p[k]
s	0	null
t	8	y
y	5	s
x	9	t
z	7	y

Poll Question #3:


Given the following table of vertices and their predecessors from Dijkstra's algorithm, what is the shortest distance between s and x?

- A. 9
- B. 8
- C. 7
- D. 5

Key (k)	d[k]	p[k]
s	0	null
t	8	y
y	5	s
x	9	t
z	7	y

Poll Question #4:

Given the following graph with **w** as the source vertex, what will the distance values be in alphabetical order after the second iteration of Dijkstra's algorithm?




- A. 0, 9, 5, ∞
- B. 0, 3, 5, 2
- C. 0, 8, 5, 7
- D. 0, 9, 5, 2

Key (k)	V[k]	dist[k]	p[k]
w	F	0	null
x	F	∞	null
y	F	∞	null
z	F	∞	null

Poll Question #4: After 1st iteration

Given the following graph with **w** as the source vertex, what will the distance values be in alphabetical order after the second iteration of Dijkstra's algorithm?




- A. 0, 9, 5, ∞
- B. 0, 3, 5, 2
- C. 0, 8, 5, 7
- D. 0, 9, 5, 2

Key (k)	V[k]	dist[k]	p[k]
w	T	0	null
x	F	9	w
y	F	5	w
z	F	∞	null

Poll Question #4: After 2nd iteration

Given the following graph with **w** as the source vertex, what will the distance values be in alphabetical order after the second iteration of Dijkstra's algorithm?




- A. 0, 9, 5, ∞
- B. 0, 3, 5, 2
- C. 0, 8, 5, 7
- D. 0, 9, 5, 2

Key (k)	V[k]	dist[k]	p[k]
w	T	0	null
x	F	8	y
y	T	5	w
z	F	7	y


Negative Edge Weight

- Dijkstra's won't find the optimal solution in a graph with negative weights.
 - Due to its Greedy Approach



Negative Weight Cycle


- A cycle with weights that sum to a negative number
- No Shortest-Path Algorithm can handle Negative Weight Cycles
- Makes sense why if you think of weights as "cost" instead of "distance"



Bellman-Ford Algorithm

- Slower ($O(|V|*|E|)$) and more complex, but more versatile than Dijkstra's
 - **Can** handle Negative Edge Weights
 - **Can** detect Negative Weight Cycles (NWC)
- Uses a *dist* map like Dijkstra's


```
for i=0 to |V|  
    for each edge (u,v) in E  
        relax the path to v in dist
```




Bellman-Ford Algorithm

- **Relaxes** ALL edges at each iteration
 - Performs *exactly* $|V|$ iterations
 - The i th iteration finds any Shortest Paths with i edges
 - Last iteration reveals presence of Negative Weight Cycles
 - The order in which you relax the edges *can* affect runtime performance

```
for i=0 to  $|V|$ 
  for each edge  $(u,v)$  in  $E$ 
    relax the path to  $v$  in  $dist$ 
```






Edges: (t,y), (s,t), (s,y)

V	d[v]	p[v]
s	0	null
t	∞	null
y	∞	null


initialize the map.
We will do $|V| = 3$
iterations, the final one
only checks for NWC



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	∞	null
y	∞	null


first iteration



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	5	s
y	∞	null


first iteration



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	5	s
y	2	s


first iteration



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	5	s
y	-5	t


second iteration



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	5	s
y	-5	t


second iteration



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	5	s
y	-5	t


second iteration



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	5	s
y	-5	t


final iteration checks
for NWC. Should be
NO updates



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	5	s
y	-5	t

final iteration checks
for NWC. Should be
NO updates



Edges: (t,y), (s,t), (s,y)

v	d[v]	p[v]
s	0	null
t	5	s
y	-5	t

final iteration checks
for NWC. Should be
NO Relaxation

Bellman Ford Pseudocode

// Step 1: Initialize graph and map for distances and predecessors

// Step 2: Relax edges repeatedly

repeat $|V| - 1$ times:

 for each edge (u,v) with weight w in edges do

 if $\text{distance}[u] + w < \text{distance}[v]$ then

$\text{distance}[v] = \text{distance}[u] + w$

$\text{predecessor}[v] = u$

Relax

// Step 3: Check for negative-weight cycles

for each edge (u,v) with weight w in edges do

 if $\text{distance}[u] + w < \text{distance}[v]$ then

 error "Graph contains a negative-weight cycle"

Poll Question #5:

Bellman-Ford is able to find the Shortest Path in any graph with negative edge weights

A. True

B. False, not if there is a NWC

Poll Question #6:

Which real-world problem would be a better fit for the Bellman-Ford Algorithm (as opposed to Dijkstra's Algorithm)?


- A. Shortest drive from Miami to Seattle
- B. Cheapest drive from Miami to Seattle
- C. Coldest drive from Miami to Seattle (temp can be negative)
- D. Hottest drive from Miami to Seattle (longest Path, different set of Algorithms)

Questions about anything?

Minimum Spanning Tree Algorithms

Minimum Spanning Tree (MST)

- MST is a subset of the edges of a **connected, weighted, and undirected** graph that:
 - connects all the vertices together
 - without any cycles (acyclic)
 - with the minimum possible total edge weight
- Two famous MST algorithms are:
 - Prim's algorithm (tracks vertices)
 - Kruskal's algorithm (tracks edges)
- Both Prim's and Kruskal's algorithms are **greedy**
- MST of a graph is **NOT** necessarily unique
 - unique only if the edge weights are unique




Poll Question #7

What is cost of the **MST** of the following Graph?

- A. 15
- B. 11
- C. 7
- D. 5

Graph:




Poll Question #7

What is cost of the **MST** of the following Graph?

- A. 15
- B. 11
- C. 7**
- D. 5


Graph:




Prim's Algorithm

- Given a **connected, weighted, and undirected** graph Prim's results in a Minimum Spanning Tree (MST)!!
- Time complexity:
 - Adjacency Matrix:
 - $O(|V|^2)$
 - Adjacency List + heap:
 - $O(|E| \log(|V|))$

Graph:




MST:




Prim's Algorithm Pseudocode

1. Create an empty set
2. Put an arbitrary node in the set
3. While your set does **not** contain all vertices of the graph
 - a. Add an adjacent vertex:
 - i. with **least** edge weight
 - ii. is **not** in the set already

Graph:



MST:




Prim's Algorithm Example

- Create a set (initially empty) of MST vertices

Set: {}

Graph:




MST:

Prim's Algorithm Example

- Pick any vertex (in this case A) and add it to your set

Graph:



Set: {A}

MST:




Prim's Algorithm Example


- While your set does not contain all vertices of the graph
 - Add an adjacent vertices:
 - with **least** edge weight
 - is **not** in the set already
- The adjacent vertices are: B, C.
B's edge is the **least** weighted, so it is added to our set.

Set: {A, B}

Graph:



MST:




Prim's Algorithm Example


- While your set does not contain all vertices of the graph
 - Add an adjacent vertices:
 - with **least** edge weight
 - is **not** in the set already
- The adjacent vertices are: C, D, and E. C's edge to B is the **least** weighted, so it is added to our set.

Set: {A, B, C}

Graph:



MST:




Prim's Algorithm Example


- While your set does not contain all vertices of the graph
 - Add an adjacent vertices:
 - with **least** edge weight
 - is **not** in the set already
- The adjacent vertices are: D, and E. D's edge to B is the **least** weighted, so it is added to our set.

Set: {A, B, C, D}

Graph:



MST:




Prim's Algorithm Example


- While your set does not contain all vertices of the graph
 - Add an adjacent vertices:
 - with **least** edge weight
 - is **not** in the set already
- The only adjacent vertice is E
E's edge to D is the **least** weighted, so it is added to our set.

Set: {A, B, C, D, E}

Graph:



MST:




Prim's Algorithm Example


- Our set includes all vertices, so we exit the while loop and our MST is complete.

Set: {A, B, C, D, E}

Graph:



MST:



Poll Question #8:

You can use Prim's Algorithm to find the MST of a graph with which of the following properties?

1. Cyclic
2. Weighted
3. Directed
4. Undirected
5. Connected

Poll Question #8:


You can use Prim's Algorithm to find the MST of a graph with which of the following properties?

1. Cyclic
2. Weighted
3. Directed
4. Undirected
5. Connected


Kruskal's Algorithm

- Given a **connected, weighted, and undirected** graph Kruskal's results in a MST
- Kruskal's forms a MST by finding the **least** weighted edges that do **NOT** form cycles (as opposed to vertices with Prim's)
- Time complexity:
 - $O(|E|\log(|E|))$, the limiting factor is sorting the edges

Graph:




MST:




Kruskal's Algorithm Pseudocode

- Create an empty set of edges
- Create a list of the edges sorted by **least** weight
- While all the vertices are **NOT** connected by the edges in our set:
 - Add the **least** weighted edge that does not form a cycle with the other edges in the set (One way to do this is with **union find**)

Graph:




MST:



Kruskal's Algorithm Example

- Create an empty set of edges

Graph:




Set: {}

MST:

Kruskal's Algorithm Example

- Create a list of the edges sorted by **least** weight

Graph:



Edges:

A-B, C-B, D-E, A-C, B-D, B-E, C-D


MST:

Set: {}

Kruskal's Algorithm Example

- A-B is the **least** weighted edge, and it does not form a cycle
 - Add it to the set

Graph:



Edges:

A-B, C-B, D-E, A-C, B-D, B-E, C-D

Set: {A-B}


MST:



Kruskal's Algorithm Example

- C-B is tied with D-E for the **least** weighted edge that does **NOT** form a cycle, so it does not matter which you pick.
- We will add C-B

Graph:




Edges:

C-B, D-E, A-C, B-D, B-E, C-D

Set: {A-B, C-B}


MST:



Kruskal's Algorithm Example

- D-E is the **least** weighted edge, and it does not form a cycle
 - Add it to the set

Graph:




Edges:

D-E, A-C, B-D, B-E, C-D

Set: {A-B, C-B, D-E}


MST:



Kruskal's Algorithm Example

- A-C is tied with B-D for the **least** weighted edge that does not form a cycle, **BUT** adding A-C would form a cycle, so we add B-D

Graph:




Edges:

A-C, B-D, B-E, C-D

Set: {A-B, C-B, D-E, B-D}

MST:



Kruskal's Algorithm Example


- All vertices have been connected, so we are done!

Edges:


B-E, C-D

Set: {A-B, C-B, D-E, B-D}

Graph:



MST:




Poll Question #9

What is the **MST** of the following Graph:

1. (A, B), (C, B), (B, D), (A, E), (E, F)
2. (A, B), (A, C), (C, D), (A, B), (A, E)
3. (A, B), (A, C), (B, D), (A, E), (E, F)

Graph:




Poll Question #9

What is the **MST** of the following Graph:

1. (A, B), (C, B), (B, D), (A, E), (E, F)
2. (A, B), (A, C), (C, D), (A, B), (A, E)
3. (A, B), (A, C), (B, D), (A, E), (E, F)


Graph:



Participation Activity

Prof Aman has lost his favorite hat!! Use Bellman-Ford to help Prof Aman find the shortest path to it (A to D).

A	B	C	D	E
0	∞	∞	∞	∞




Note that this is a digraph.

Use edge order: (B, E), (D, B), (B, D), (A, B), (A, C), (D, C), (B, C), (E, D)

Participation Activity

Prof Aman has lost his favorite hat!! Use Bellman-Ford to help Prof Aman find the shortest path to it (A to D).

	A	B	C	D	E
A	0	∞	∞	∞	∞
B	0	-1	∞	∞	∞
C	0	-1	4	∞	∞
D	0	-1	2	∞	∞
E					




Iteration 1: (B, E), (D, B), (B, D), (A, B), (A, C), (D, C), (B, C), (E, D)

Participation Activity

Prof Aman has lost his favorite hat!! Use Bellman-Ford to help Prof Aman find the shortest path to it (A to D).

	A	B	C	D	E
0	∞	∞	∞	∞	∞
0	-1	∞	∞	∞	∞
0	-1	4	∞	∞	∞
0	-1	2	∞	∞	∞
0	-1	2	∞	1	1
0	-1	2	1	1	1
0	-1	2	-2	1	1



Iteration 2: (B, E), (D, B), (B, D), (A, B), (A, C), (D, C), (B, C), (E, D). **Iterations 3 and 4** won't update table.

Bellman Ford Walkthrough Video

<https://www.youtube.com/watch?v=obWXjtg0L64>



The image shows a YouTube video player interface. The video title is "Bellman-Ford in 5 minutes – Step by step example". The channel name is "Michael Sambol" with a verified badge and "103K subscribers". The video has 17K likes. The player controls show a progress bar at 0:00 / 5:09, a play button, a volume icon, a closed captions icon, a settings icon, a full screen icon, and a download icon. The video content itself is a white background with the text "Bellman-Ford Algorithm" written in a cursive, handwritten style.

Bellman-Ford
Algorithm

0:00 / 5:09

Bellman-Ford in 5 minutes – Step by step example

Michael Sambol ✓
103K subscribers

Subscribe

17K

Share

Download

Extra Resources

[Min Heap Implementation of Dijkstra's Algorithm](#)

[Dijkstra's vs Bellman Ford](#)

GeeksforGeeks graph type overview:

<https://www.geeksforgeeks.org/graph-types-and-applications/>

Shortest path algorithms: AFoolsPath


<https://towardsdatascience.com/algorithm-shortest-paths-1d8fa3f50769>

<https://www.hackerearth.com/practice/algorithms/graphs/shortest-path-algorithms/tutorial/>

Reminder:

You have an HonorLock
quiz due this Friday by
11:59pm

Detecting Cycles with the Union-Find Data Structure



MAKE-SET(x)

UNION(x, y)

FIND-SET(x)

Solution to collaborative question

Using disjoint sets:

```
MST-KRUSKAL( $G, w$ )
1  $A = \emptyset$ 
2 for each vertex  $v$  of  $V$ 
3   MAKE-SET( $v$ )
4 sort the edges of  $E$  into nondecreasing order by weight  $w$ 
5 for each edge  $(u, v) \in E$ , taken in nondecreasing order by weight
6   if FIND-SET( $u$ )  $\neq$  FIND-SET( $v$ )
7     then  $A \leftarrow A \cup \{(u, v)\}$ 
8     UNION( $u, v$ )
9 return  $A$ 
```

Extra Material